

Title: N-gram

URL: <https://en.wikipedia.org/wiki/N-gram>

PageID: 986182

Categories: Category:Computational linguistics, Category:Corpus linguistics, Category:Language modeling, Category:Natural language processing, Category:Probabilistic models, Category:Speech recognition

Source: Wikipedia (CC BY-SA 4.0).

An n -gram is a sequence of n adjacent symbols in a particular order. [1] The symbols may be n adjacent letters (including punctuation marks and blanks), syllables , or rarely whole words found in a language dataset; or adjacent phonemes extracted from a speech-recording dataset, or adjacent base pairs extracted from a genome. They are collected from a text corpus or speech corpus .

If Latin numerical prefixes are used, then n -gram of size 1 is called a "unigram", size 2 a " bigram " (or, less commonly, a "digram") etc. If, instead of the Latin ones, the English cardinal numbers are furtherly used, then they are called "four-gram", "five-gram", etc. Similarly, using Greek numerical prefixes such as "monomer", "dimer", "trimer", "tetramer", "pentamer", etc., or English cardinal numbers, "one-mer", "two-mer", "three-mer", etc. are used in computational biology, for polymers or oligomers of a known size, called k -mers . When the items are words, n -grams may also be called shingles . [2]

In the context of natural language processing (NLP), the use of n -grams allows bag-of-words models to capture information such as word order, which would not be possible in the traditional bag of words setting.

Examples

In 1951, Shannon [3] discussed n -gram models of English. For example:

3-gram character model (random draw based on the probabilities of each trigram): in no ist lat whey cratict froure birs grocid pondenome of demonstures of the retagin is regiactona of cre

2-gram word model (random draw of words taking into account their transition probabilities): the head and in frontal attack on an english writer that the character of this point is therefore another method for the letters that the time of who ever told the problem for an unexpected

Figure 1 shows several example sequences and the corresponding 1-gram, 2-gram and 3-gram sequences.

Here are further examples; these are word-level 3-grams and 4-grams (and counts of the number of times they appeared) from the Google n -gram corpus. [4]

3-grams

ceramics collectables collectibles (55)

ceramics collectables fine (130)

ceramics collected by (52)

ceramics collectible pottery (50)

ceramics collectibles cooking (45)

4-grams

serve as the incoming (92)

serve as the incubator (99)

serve as the independent (794)

serve as the index (223)

serve as the indication (72)

serve as the indicator (120)

References

Further reading

Manning, Christopher D.; Schütze, Hinrich; Foundations of Statistical Natural Language Processing , MIT Press: 1999, ISBN 0-262-13360-1

White, Owen; Dunning, Ted; Sutton, Granger; Adams, Mark; Venter, J. Craig; Fields, Chris (1993). "A quality control algorithm for dna sequencing projects" . Nucleic Acids Research . 21 (16): 3829–3838. doi : 10.1093/nar/21.16.3829 . PMC 309901 . PMID 8367301 .

Damerau, Frederick J.; Markov Models and Linguistic Theory , Mouton, The Hague, 1971

Figuerola, Alejandro; Atkinson, John (2012). "Contextual Language Models For Ranking Answers To Natural Language Definition Questions". Computational Intelligence . 28 (4): 528– 548. doi : 10.1111/j.1467-8640.2012.00426.x .

Brocardo, Marcelo Luiz; Traore, Issa; Saad, Sherif; Woungang, Isaac (2013). "Authorship verification for short messages using stylometry". 2013 International Conference on Computer, Information and Telecommunication Systems (CITS) . pp. 1– 6. doi : 10.1109/CITS.2013.6705711 . ISBN 978-1-4799-0168-5 .

See also

Google Books Ngram Viewer

External links

Ngram Extractor: Gives weight of n -gram based on their frequency.

Google's Google Books n -gram viewer and Web n -grams database (September 2006)

STATOPERATOR N-grams Project Weighted n -gram viewer for every domain in Alexa Top 1M
1,000,000 most frequent 2,3,4,5-grams from the 425 million word Corpus of Contemporary American English

Peachnote's music ngram viewer

Stochastic Language Models (n -Gram) Specification (W3C)

Michael Collins's notes on n -Gram Language Models

OpenRefine: Clustering In Depth

v

t

e

AI-complete

Bag-of-words

n -gram Bigram Trigram

Bigram

Trigram

Computational linguistics

Natural language understanding

Stop words

Text processing
Argument mining
Collocation extraction
Concept mining
Coreference resolution
Deep linguistic processing
Distant reading
Information extraction
Named-entity recognition
Ontology learning
Parsing Semantic parsing Syntactic parsing
Semantic parsing
Syntactic parsing
Part-of-speech tagging
Semantic analysis
Semantic role labeling
Semantic decomposition
Semantic similarity
Sentiment analysis
Terminology extraction
Text mining
Textual entailment
Truecasing
Word-sense disambiguation
Word-sense induction
Compound-term processing
Lemmatisation
Lexical analysis
Text chunking
Stemming
Sentence segmentation
Word segmentation
Multi-document summarization
Sentence extraction
Text simplification
Computer-assisted
Example-based
Rule-based

Statistical
Transfer-based
Neural
BERT
Document-term matrix
Explicit semantic analysis
fastText
GloVe
Language model (large)
Latent semantic analysis
Seq2seq
Word embedding
Word2vec
Corpus linguistics
Lexical resource
Linguistic Linked Open Data
Machine-readable dictionary
Parallel text
PropBank
Semantic network
Simple Knowledge Organization System
Speech corpus
Text corpus
Thesaurus (information retrieval)
Treebank
Universal Dependencies
BabelNet
Bank of English
DBpedia
FrameNet
Google Ngram Viewer
UBY
WordNet
Wikidata
Speech recognition
Speech segmentation
Speech synthesis
Natural language generation

Optical character recognition
Document classification
Latent Dirichlet allocation
Pachinko allocation
Automated essay scoring
Concordancer
Grammar checker
Predictive text
Pronunciation assessment
Spell checker
Chatbot
Interactive fiction
Question answering
Virtual assistant
Voice user interface
Formal semantics
Hallucination
Natural Language Toolkit
spaCy